


No tested model preserves allocation at full construction



Every constructive endpoint is R3 (magnitude-100, same problem instances). GPT entries are conditional on commitment: GPT-5.4-mini realizes 34.7% of write-budget slots, Sol 73.6% (front-loaded into each session's opening problems). Qwen points: one matched A2 panel (24 held-out streams, q8_0); base is eager in both frames, the RL-trained policy reserves abstractly but not during construction.